

ASSIGNMENT

Problem Statement

- Install Logstash. This can be on a PC/laptop or in a virtual machine.
- Create a Logstash configuration file. The Logstash configuration file should do the following:
 - Allow for ingestion of the **Security Log Data** below via a “Logstash Input Plugin”, your choice.
 - Allow for output of data via a “Logstash Output Plugin”, your choice.
 - Utilizing any number of “Logstash Filter Plugins”, transform data from the **Security Log Data** from its original format to a normalized structure described below.
- Design your solution so that a different piece of **Security Log Data** structured the same, but containing different data in the log message, could be ingested and achieve a similar output based on the data contained within the log.
- Send the **Security Log Data** example to Logstash and capture the output from Logstash in some form (text, screenshot, etc). Include this in the Git repo.
- Save the Logstash configuration created above in a public Git repo.
- Write a short description of how data is ingested and output from Logstash based on your solution. Include this write-up in the Git repo.

Security Log Data:

<14>1 2016-12-25T09:03:52.754646-06:00 acmeinchost1 antivirus 2496 - - alertname="Virus Found" computername="acmeincpc42" computerip="10.58.194.142" severity="1"

Output as JSON formatted data containing (but not exclusively) the following fields:

description: "Virus Found"

hostname: " acmeincpc42"

source_ip: "10.58.194.142"

severity: "High"

Solution

1. Installed Elasticsearch, Kibana and Logstash and configured to connect through local server
2. Verified the working of ELK

- I used the **stdin** and **stdout** as input and output plugin to do **console testing** wherein, I am ingesting data (a sample log) from the terminal and seeing the output on the terminal as well.
- Initially, I created the config file by directly parsing the fields from the log in the **grok** pattern. Attached below is the config file and the output:



Config File.txt

Output

```
Windows PowerShell
[2025-04-05T23:29:58,936][WARN ][logstash.filters.grok      ][main] ECS v8 support is a preview of the unreleased ECS v8, and uses the
v1 patterns. When Version 8 of the Elastic Common Schema becomes available, this plugin will need to be updated
[2025-04-05T23:29:58,938][INFO ][logstash.outputs.elasticsearch][main] Using a default mapping template {es_version=>8, :ecs_compati
bility=>v8}
[2025-04-05T23:29:59,096][INFO ][logstash.javapipeline     ][main] Starting pipeline {pipeline_id=>"main", "pipeline.workers"=>8, "pi
pline.batch.size"=>125, "pipeline.batch.delay"=>50, "pipeline.max_inflight"=>1000, "pipeline.sources"=>["D:/Elastic Stack/Logstash/l
ogstash-8.17.4/config/logstash.conf"], :thread=>"#<Thread:0x7e3f7da3 D:/Elastic Stack/Logstash/logstash-8.17.4/logstash-core/lib/logs
tash/java_pipeline.rb:138 run>"}
[2025-04-05T23:30:00,435][INFO ][logstash.javapipeline     ][main] Pipeline Java execution initialization time {"seconds"=>1.34}
[2025-04-05T23:30:00,471][INFO ][logstash.javapipeline     ][main] Pipeline started {"pipeline.id"=>"main"}
The stdin plugin is now waiting for input:
[2025-04-05T23:30:00,485][INFO ][logstash.agent             ][main] Pipelines running {:count=>1, :running_pipelines=>[:main], :non_running_p
ipelines=>[]}
1 2016-12-25T09:03:52.754646-06:00 acmeinhost1 antivirus 2496 - - alertname="Virus Found" computername="acmeincpc42" computerip="10.
58.194.142" severity="1"
{
  "event_type" => "antivirus",
  "description" => "Virus Found",
  "device_name" => "acmeincpc42",
  "severity"    => "High",
  "hostname"    => "acmeinhost1",
  "source_ip"   => "10.58.194.142",
  "@timestamp"  => 2016-12-25T15:03:52.754Z,
  "pid"         => "2496"
}
```

- Further **fine-tuning** the config file, I used the **kv** filter to segregate the remaining part of the message which had a key and a value. But the output had a backtracking due to **severity** value giving **string** values when it is configured to support **long-int** values



Config_File_With_KV
P_Technique.txt

Output

```
Windows PowerShell
[2025-04-05T23:48:05,742][WARN ][logstash.filters.grok      ][main] ECS v8 support is a preview of the unreleased ECS v8, and uses the
v1 patterns. When Version 8 of the Elastic Common Schema becomes available, this plugin will need to be updated
[2025-04-05T23:48:05,748][INFO ][logstash.outputs.elasticsearch][main] Using a default mapping template {es_version=>8, :ecs_compati
bility=>v8}
[2025-04-05T23:48:06,079][INFO ][logstash.javapipeline     ][main] Starting pipeline {pipeline_id=>"main", "pipeline.workers"=>8, "pi
pline.batch.size"=>125, "pipeline.batch.delay"=>50, "pipeline.max_inflight"=>1000, "pipeline.sources"=>["D:/Elastic Stack/Logstash/l
ogstash-8.17.4/config/logstash.conf"], :thread=>"#<Thread:0xde792bf D:/Elastic Stack/Logstash/logstash-8.17.4/logstash-core/lib/logst
ash/java_pipeline.rb:138 run>"}
[2025-04-05T23:48:08,372][INFO ][logstash.javapipeline     ][main] Pipeline Java execution initialization time {"seconds"=>2.29}
[2025-04-05T23:48:08,518][INFO ][logstash.javapipeline     ][main] Pipeline started {"pipeline.id"=>"main"}
The stdin plugin is now waiting for input:
[2025-04-05T23:48:08,559][INFO ][logstash.agent             ][main] Pipelines running {:count=>1, :running_pipelines=>[:main], :non_running_p
ipelines=>[]}
1 2016-12-25T09:03:52.754646-06:00 acmeinhost1 antivirus 2496 - - alertname="Virus Found" computername="acmeincpc42" computerip="10.
58.194.142" severity="1"
{
  "source_ip" => "10.58.194.142",
  "pid"       => "2496",
  "event_type" => "antivirus",
  "description" => "Virus Found",
  "device_name" => "acmeincpc42",
  "hostname"    => "acmeinhost1",
  "severity"    => "High",
  "@timestamp"  => 2016-12-25T15:03:52.754Z
}
```

In both the above configs, I used the **rubydebug** codec to parse the output data

6. Analyzing further, I see that using a different label for **severity** field would be an ideal approach as I observed that field mappings were blocked.

Hence, I modified the config file to parse **severity** in an additional field **device_severity** and mapped the **string** value to this field.

Input Plugin

stdin

Filters Used

grok – To match the header part of the log

kv – To split the key and value from the remaining part of the log

mutate + **gsub** – To remove any double quotes in the severity value which has an integer

translate – For converting integer severity values to strings with “Low”, “Medium” and “High” severity levels

mutate + **rename** – Renamed the field names given in the **grok** pattern to match to match the problem statement

mutate + **remove_fields** – Removed some of the fields from the **stdout** output JSON

Output Plugin

stdout

Elasticsearch Index and Host

hosts => ["http://localhost:9200"]

index => "test.logstash"

Attached below is the config file and the output:



logstash.conf

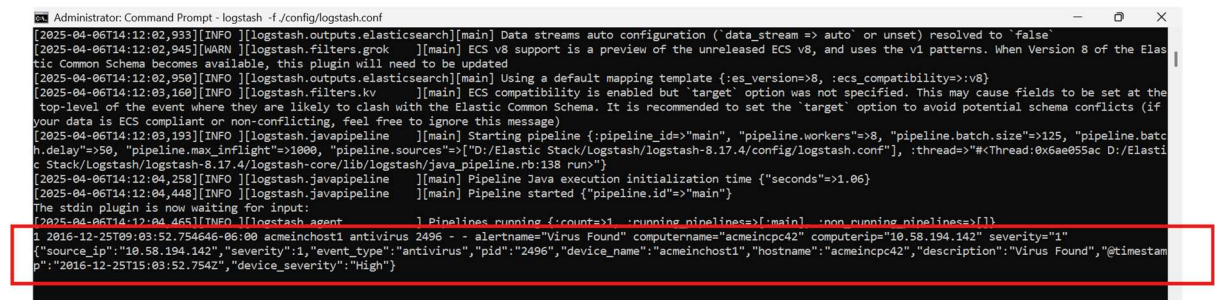
Input

```
D:\Elastic Stack\logstash\logstash-8.17.4\bin>logstash -f .\config\logstash.conf
Using bundled JDK: D:\Elastic Stack\logstash\logstash-8.17.4\jdk\bin\java.exe
Sending Logstash logs to D:\Elastic Stack\logstash\logstash-8.17.4\logs which is now configured via log4j2.properties
[2025-04-06T12:50:51,901][WARN ][logstash.runner] NOTICE: Running Logstash as a superuser is strongly discouraged as it poses a security risk. Set 'allow_superuser' to false for better security.
[2025-04-06T12:50:51,907][INFO ][logstash.runner] Log4j configuration path used is: D:\Elastic Stack\logstash\logstash-8.17.4\config\log4j2.properties
[2025-04-06T12:50:51,907][INFO ][logstash.runner] Starting Logstash {"logstash.version":"8.17.4", "jruby.version":"jruby 9.4.9.0 (3.1.4) 2024-11-04 547c6b150e 0
penDK 64-Bit Server VM 21.0.6+7-LTS on 21.0.6+7-LTS +indy +jit [x86_64-mswin32]"}
[2025-04-06T12:50:51,907][INFO ][logstash.runner] JVM bootstrap flags: [-Xms1g, -Xmx1g, -Djava.awt.headless=true, -Dfile.encoding=UTF-8, -Djruby.compile.invokedyn
amic=true, -XX:+HeapDumpOnOutOfMemoryError, -Djava.security.egd=file:/dev/urandom, -Dlog4j2.isThreadContextMapInheritable=true, -Dlogstash.jackson.stream.read.constraints.m
ax-string-length=200000000, -Dlogstash.jackson.stream.read.constraints.max-number-length=10000, -Djruby.regexp.interruptible=true, -Djdk.io.File.enableADS=true, --add-expor
ts=jdk.compiler/com.sun.tools.javac.api=ALL-UNNAMED, --add-exports=jdk.compiler/com.sun.tools.javac.file=ALL-UNNAMED, --add-exports=jdk.compiler/com.sun.tools.javac.parser=
ALL-UNNAMED, --add-exports=jdk.compiler/com.sun.tools.javac.tree=ALL-UNNAMED, --add-exports=jdk.compiler/com.sun.tools.javac.util=ALL-UNNAMED, --add-opens=java.base/java.se
curity=ALL-UNNAMED, --add-opens=java.base/java.io=ALL-UNNAMED, --add-opens=java.base/java.nio.channels=ALL-UNNAMED, --add-opens=java.base/sun.nio.ch=ALL-UNNAMED, --add-open
s=java.management/sun.management=ALL-UNNAMED, -Dio.netty allocator.maxOrder=11]
[2025-04-06T12:50:51,965][INFO ][org.logstash.jackson.StreamReadConstraintsUtil] Jackson default value override 'logstash.jackson.stream.read.constraints.max-string-length'
configured to '200000000'
[2025-04-06T12:50:51,965][INFO ][org.logstash.jackson.StreamReadConstraintsUtil] Jackson default value override 'logstash.jackson.stream.read.constraints.max-number-length'
configured to '10000'
[2025-04-06T12:50:52,015][WARN ][logstash.config.source.multilocal] Ignoring the 'pipelines.yml' file because modules or command line options are specified
[2025-04-06T12:50:52,941][INFO ][logstash.agent] Successfully started Logstash API endpoint {"port=>9600, :ssl_enabled=>false}
[2025-04-06T12:50:53,501][INFO ][org.reflections.Reflections] Reflections took 96 ms to scan 1 urls, producing 152 keys and 530 values
[2025-04-06T12:50:55,531][WARN ][logstash.filters.translate] You are using a deprecated config setting "destination" set in translate. Deprecated settings will continue to
```

Output

```
{
  "source_ip": "10.58.194.142",
  "severity": 1,
  "event_type": "antivirus",
  "pid": "2496",
  "device_name": "acmeinchost1",
  "hostname": "acmeincpc42",
  "description": "Virus Found",
  "@timestamp": "2016-12-25T15:03:52.754Z",
  "device_severity": "High"
}
```

Final Output from Terminal



```
Administrator: Command Prompt - logstash -f ./config/logstash.conf
[2025-04-06T14:12:02.933][INFO ][logstash.outputs.elasticsearch][main] Data streams auto configuration ('data_stream' => auto) or unset) resolved to 'false'
[2025-04-06T14:12:02.945][WARN ][logstash.filters.grok      ][main] ECS v8 support is a preview of the unreleased ECS v8, and uses the v1 patterns. When Version 8 of the Elastic Common Schema becomes available, this plugin will need to be updated
[2025-04-06T14:12:02.950][INFO ][logstash.outputs.elasticsearch][main] Using a default mapping template {'es_version':8, 'ecs_compatibility':v8}
[2025-04-06T14:12:03.160][INFO ][logstash.filters.kv      ][main] ECS compatibility is enabled but 'target' option was not specified. This may cause fields to be set at the top-level of the event where they are likely to clash with the Elastic Common Schema. It is recommended to set the 'target' option to avoid potential schema conflicts (if your data is ECS compliant or non-conflicting, feel free to ignore this message)
[2025-04-06T14:12:03.193][INFO ][logstash.javapipeline   ][main] Starting pipeline {:pipeline_id=>"main", "pipeline.workers">8, "pipeline.batch.size">125, "pipeline.batch.delay">50, "pipeline.max_inflight">1000, "pipeline.sources">["D:/Elastic Stack/Logstash/logstash-8.17.4/logstash-core/lib/logstash/java_pipeline.rb:138 run"]}
[2025-04-06T14:12:04.258][INFO ][logstash.javapipeline   ][main] Pipeline java execution initialization time {"seconds">1.06}
[2025-04-06T14:12:04.448][INFO ][logstash.javapipeline   ][main] Pipeline started {"pipeline.id">"main"}
The stdin plugin is now waiting for input:
[2025-04-06T14:12:04.465][INFO ][logstash.agent         ][main] Pipelines running {:count=>1, :running_pipelines=>[main], :non_running_pipelines=>[]}
1 2016-12-25T09:03:52.754646-06:00 acmeinchost1 antivirus 2496 - - alertname="Virus Found" computername="acmeincpc42" computerip="10.58.194.142" severity="1"
{"source_ip":"10.58.194.142","severity":1,"event_type":"antivirus","pid":"2496","device_name":"acmeinchost1","hostname":"acmeincpc42","description":"Virus Found","@timestamp":"2016-12-25T15:03:52.754Z","device_severity":"High"}
```

In the above output, I mapped **string** values of **severity** to **device_severity** and used **JSON** as the **codec**.

The remaining attributes are mapped as per what is expected in the output.

Hence, the final version of the config file that I came up with, is attached below:



logstash.conf